Second Life Commerce: AI Powered Returns

ReLoop

The intelligence layer for Amazon's return pipeline.

Grade at the doorstep, decide the item's next life **before it moves.**

The smart layer that sits on top of Amazon's returns, so far less gets wasted.

START WITH SOMETHING FAMILIAR

What actually happens after you click “Return”?

You hand the package back and forget about it. Behind the scenes, it goes on a surprisingly long trip:

1. **Picked up** › a driver collects it from your door
2. **Local depot** › it stops at a nearby station just passing through
3. **Sorting center** › trucked to a bigger hub to be batched up
4. **Returns center** › hauled hundreds of km away often 500km+
5. **Waits in line** › sits for days or weeks before anyone looks at it
6. **Finally graded** only NOW does a person decide what it's worth

The catch: the single most important decision is that “what should we do with this item?" happens **dead last**, after all the driving is already paid for and weeks f value have quietly leaked away.

WHY IT MATTERS

That long trip is expensive and most items lose value

Returns are enormous, and today's process wrings very little value back out of them.

1.2–1.5 billion

packages returned every year

₹30,000-50,000 Crore

what returns cost the industry, yearly

only 10–20%

of returns ever make it back onto a shelf to be resold

the rest

mostly sold off cheap to liquidators or are thrown away

Every wasted item is money burned twice: once on the trucking, and again on the value that could have been recovered if we'd acted sooner.

THE IDEA

So… what if we decided at your front door instead?

The whole problem is that we decide **after** the item has travelled. ReLoop flips it: look at the item the moment it's returned, and pick its best next step **before it moves an inch.**

TODAY

- Ship it far away first
- Let it sit and lose value
- Then decide what is it's worth
- Usually too late to do anything clever

WITH RELOOP

- AI decides the grade at the doorstep, instantly
- Valuable items are kept local and rerouted to nearby buyers
- Skip the pointless long-distance trip to the warehouse
- Only the items that truly need it travel, so it is purely additive

THE WHOLE THING IN THREE STEPS

Snap a photo → the AI decides → the item finds its best home

1

Look at it

You take a few photos at the door. An AI reads the item's condition like a really fast, really consistent inspector.

2

Think it through

A decision engine intercepts the return before it moves and reroutes it to the highest-value local path, instead of defaulting to a warehouse 200km away.

3

Send it the right way

The item goes straight to its best next home. Two quick human checks along the way make sure the AI got it right.

STEP 1: THE EYES

An AI reads the item's condition from photos

- **Trained to judge condition.** It looks at the photos and figures out how good a shape the item is in: brand-new, lightly used, worn, or beyond saving.
- **It spots the details.** Scratched screen? Torn box? It notices, and lists them into the product card.
- **And here's the important part: I**t's honest about how sure it is. Instead of forcing a single answer, it outputs a confidence score so the system knows when to trust the grade and when to flag it for a closer look.

Why "how sure" matters

A confident wrong guess is dangerous, imagine selling a scratched phone to someone as "like new." **By keeping track of its own doubt, the engine can play it safe when the photos aren't clear.**

```
what the AI hands over:

"70% like-new, 20% lightly used,
 10% worn — fairly confident.
 Note: box looks a bit scuffed."
```

What's the best thing to do with this item, right now?

For every item, it weighs up the full picture and not just what it's worth, but what it costs to get it there, who actually wants it nearby, and how fast the clock is ticking and then picks the best path.

What it's worth

the resale value for the condition it's in

The costs to get there

trucking, handling, repairs subtracted honestly

Who wants it nearby

how many local buyers are looking for this right now

Time

things lose value while they sit and that counts as a cost too

The planet

carbon saved by not trucking it far is worth money here

Cost of a mistake

the penalty if this route fails and needs correcting

No black box. It's plain arithmetic and clear rules and you can read exactly why it chose what it chose. The AI only writes the one-line explanation; the **rules do the deciding.**

SEE IT IN ACTION

A real example: a returned wireless speaker

Priya ordered a JBL Flip 6 for ₹6,499. It arrived fine but she just didn't like the sound. Box unopened, return initiated.
ReLoop's AI grades it in 3 seconds: **Grade A, confidence 0.91.** The engine then runs every option simultaneously

```
the engine's math, in rupees

Best option chosen:  put it back on sale

Resell Nearby +₹5,720   ← winner
Back on the shelf +₹5,850
Repair & resell       +₹4,100
Sell in a graded lot    +₹2,300
Donate                  +₹182

Ship it far away        −₹465   ← today's
                                  default
```

Default Idea

Today's default is shipping it far away and it is the **worst** option here (it loses money). Every smarter choice is money we get to keep.

THE OBVIOUS WORRY

“But what if the AI gets it wrong?”

Fair question. Three things make a wrong call cheap and rare and not a disaster.

1

When unsure, it plays safe

The less confident the AI is, the more cautious the option it's allowed to pick. Below a confidence threshold, it skips local routing entirely and defaults to the warehouse.

2

Two humans double-check

A 30-second look by the pickup driver, then a 10-minute check at the local station (Amazon FBA). If a person disagrees, the plan updates on the spot.

3

Worst case = today

If nothing local makes sense, it does exactly what happens today and sends it to the warehouse. ReLoop can only do better, never worse.

THE TRICKY CASES

“But what about the weird ones?”

Every odd case lands on one of two guarantees: play it safe, or fall back to exactly what Amazon does today. So we never do worse.

The tricky case	The worry	How ReLoop handles it
The AI grades it wrong	A bad call reaches a buyer	Caught by two human checks (driver + station) before any buyer sees it — worst case is one shelf move
Too good to bother (sealed, like-new)	Wasted on a cheap channel	Straight back on the shelf — skips the pipeline, sells before it loses value
The AI isn't sure (blurry photos)	An over-confident wrong guess	The less sure it is, the safer it plays — low confidence falls back to today's warehouse flow
Packaging / accessories missing	Can't resell as new	Repacked at the local station; a missing ₹300 cable becomes a repair worth + ₹1,500
Nobody nearby wants it	The item just sits unsold	Demand was priced in up front; an unsold listing auto-cascades to donate / recycle
Customer hides damage / fraud	Refunding a lie	Refund is held until a human verifies; photo-reuse + capture checks flag fakes
Every option loses money (cheap item)	Costs more to process than it's worth	“Keep it, we'll refund you” — no pickup at all; trusted customers only, never with a fraud signal
Counterfeit / hazmat / recalled	Legal & safety risk	Hard safety rules decide first, never overridden by money — routed to seller or certified disposal

THE TRUST QUESTION: WOULD A STRANGER BUY IT?

THE TRUST QUESTION: WILL THE SELLER LET GO?

Every item gets a health record so the seller knows exactly what they're sending

Sellers hesitate to accept local routing because they don't know what their item is actually worth in its current state. The Health Card removes that hesitation.

- **A Health Card is generated the moment grading happens** condition, photos, every defect spotted, confidence score.
- **It tells the seller exactly what they have,** not a guess, a verified snapshot so they can make the local routing decision with confidence.

It travels with the item to the hub and the buyer, so everyone in the chain sees the same verified truth.

PRODUCT HEALTH CARD

✓ Condition: like-new (verified)

✓ Checked at pickup

✓ Checked at the local station

✓ Genuine — serial matches

• Photos included

• Owner 2 of 2 · CO_2 saved so far

IT DOESN'T STOP AT “LIST IT”

An agent keeps working until the item sells

- It watches each listing and adjusts the price as time runs out and demand shifts something a warehouse queue can never do.
- **It finds the right buyers** by tapping into who nearby already searched for, wishlisted, or bought similar things.
- **It knows when to give up gracefully.** If resale truly isn't working, it doesn't let the item rot it hands it back to the engine, which donates or recycles it instead.

Where a buyer sees it

Verified open-box near you which is checked, guaranteed, delivered today.

```
A returns warehouse can't
negotiate a price or chase a
buyer. This little agent does —
every single day.
```

FOLLOW THE MONEY

Why this actually saves and makes money

Skip the long trip

For local items, we delete the sorting hop, the 500km+ haul, and the weeks of waiting. That's pure saved cost.

Sell before value drops

Electronics and fashion lose value fast. Selling in days instead of weeks keeps more of the price.

Match before it moves

A buyer is found during the pickup window so the item travels once, direct to its next owner. No warehouse stop, no second shipment.

Better liquidation

Clearly-labelled, graded boxes sell for more than “mystery pallets,” because buyers can see what's inside.

WHY THIS IS AMAZON'S TO WIN

Amazon already owns the rails and we add the missing piece

- **The trucks, stations, and resale stores already exist.** ReLoop doesn't rebuild any of that, it just adds a brain at the very start.
- **The trust is already there.** Every item is Amazon-checked and guaranteed and not a stranger's promise. This is why it's not just another used-goods app.
- **It's a safe bet.** Worst case, it does exactly what happens today. Every good early decision on top of that is pure upside.

In one line: look at the item at your door, decide its best next home before it moves and waste far less.

UNDER THE HOOD — THE EYES

Our own grading model — not an API call

- **We trained our own eyes.** A DINOv2 vision model — not a generic API — so grading can run right at the doorstep, with no cloud round-trip per return.
- **One photo, four answers at once.** The grade, how sure it is, the exact defects and their severity, and an overall damage score.
- **Taught on real + fake damage.** Clean Amazon catalog photos, computer-made damage with known labels, and real-world defect datasets.
- **Honest, not just confident.** It's calibrated — “80% sure” really means 80%. We cut its over-confidence by about 75%.

What the model hands back

```
grade:        70% A · 22% B · 8% C
confidence:   0.91   (calibrated)
defects:      scratched screen · torn box
damage:       0.12 overall
```

Two things an API can't do

Reference check: matches the returned item to its original catalog photo + reads the serial — catches swaps & fakes.

Flywheel: every human check at the hub is a free training label — it gets better just by doing the job.

A living price — and buyers found in priority order

The price never sits still

```
price = base × condition × demand
        × urgency × category
```

- **Real local demand,** a regional demand index refreshed hourly from what buyers actually search, view and buy nearby.
- **Racing the clock,** urgency climbs as the pickup window closes, so the price adapts on its own.
- **A pricing agent,** XGBoost + a bandit nudges the listed price toward what similar items are truly selling for — a warehouse queue can't.

We notify buyers in order — not all at once

```
score = .30 near + .35 intent + .20 fit + .15 recency
```

- **Intent leads,** who already searched, wishlisted, or bought similar things nearby ranks highest.
- **Notify #1 first,** if they don't bite in the window, it cascades to #2, #3… automatically — a job that runs every 30 min and survives restarts.
- **Fair matches only,** within 10 km, same city, and inside each buyer's budget and condition floor.

One deterministic engine, re-run at every checkpoint

For every path: value (across every possible grade × how fast it sells) − (shipping + handling + cost of being wrong) − carbon → pick the biggest number.

It thinks in probabilities

Not one guess — a whole distribution. Restock is punished hard for a wrong “A”; donating barely cares.

It prices time

Weeks of warehouse waiting is a real cost line, so “sell it locally now” can beat “ship it and wait.”

It prices being wrong

The cost of a mistake on each path sets how confident the AI must be before that path is even allowed.

- **Safety first.** A hard rule ladder (counterfeit, hazmat, recall) runs before any money math and is never overridden.
- **Glass box.** It's plain arithmetic and clear rules — the AI only writes the one-line reason. And it re-runs at pickup and at the hub as the evidence improves.

BUILT TO SCALE ON AWS

How this runs at Amazon scale

The piece	Rides on AWS
The grading model	SageMaker endpoints + Greengrass at the edge for true doorstep inference; Bedrock vision for narration & fallback
Photos & Health Cards	S3 — with Rekognition redacting faces and addresses right at the upload boundary
The decision engine	Stateless Lambda — a pure, replayable function that scales out flat
The return's journey	Step Functions + EventBridge fire the engine at each checkpoint; item state in DynamoDB
Pricing & buyer matching	Hourly batch demand index, live price on read; EventBridge Scheduler + geo DynamoDB drive the cascade, SNS / Pinpoint send the alerts
When the AI's unsure	Amazon A2I human-review queues; CloudWatch + SageMaker Model Monitor watch for drift
Getting smarter	Hub verdicts stream to a Feature Store / S3 data lake → periodic retraining

Every piece is stateless or event-driven — it rides Amazon's existing trucks, stations and stores. No new hardware, no new buildings.

ReLoop

Grade at the doorstep.

Decide before the item moves.

“The landfill is a design flaw.”

reloop-woad.vercel.app